

# Deep Reinforcement Learning for Autonomous Systems

---

Designing a control system to exploit model-free deep reinforcement learning algorithms to solve a real-world autonomous driving task of a small robot.

**Candidate:** Piero Macaluso

**Supervisors:** Prof. Pietro Michiardi  
Prof. Elena Baralis

EURECOM, France  
Politecnico di Torino, Italy

March 12, 2020



**POLITECNICO  
DI TORINO**

This work of this thesis was developed at EURECOM (Sophia Antipolis, France)  
in collaboration with

Prof. Pietro Michiardi (EURECOM)

Prof. Elena Baralis (Politecnico di Torino)


# Table of contents

1. Background
2. Design of the control system
3. Experimental methodology and results
4. Conclusions and future work


# Background

---

# State-of-the-art Autonomous Driving Systems

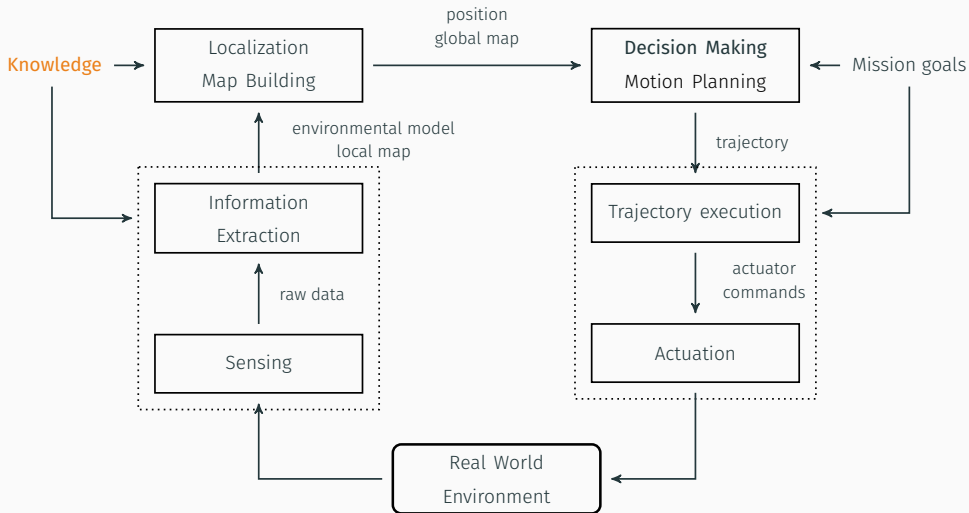


# State-of-the-art Autonomous Driving Systems



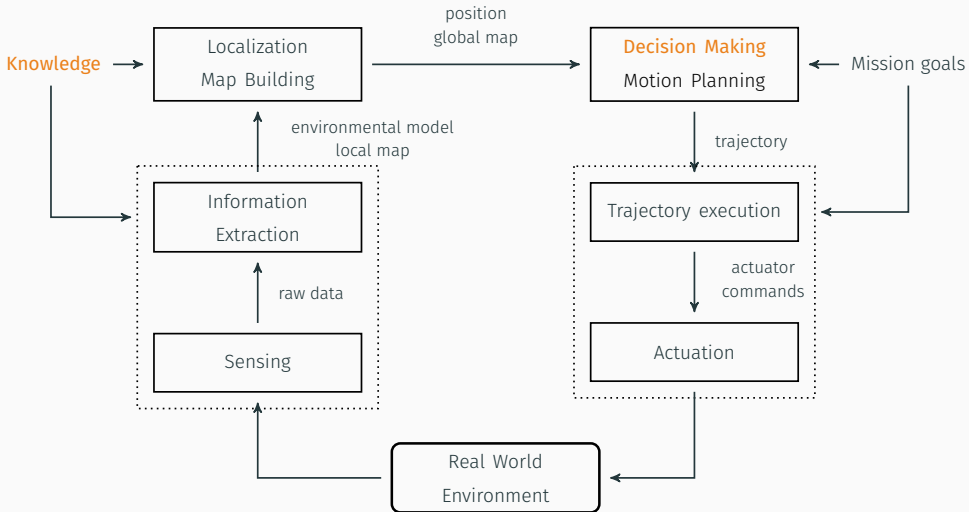
Deep Learning and Machine Learning are mainly exploited in  
object detection and recognition.

# State-of-the-art Autonomous Driving Systems



Pavone, *Veicoli a guida autonoma: a che punto siamo e cosa ci aspetta?*

# State-of-the-art Autonomous Driving Systems




## Pavone, Veicoli a guida autonoma: a che punto siamo e cosa ci aspetta?



# Reinforcement Learning

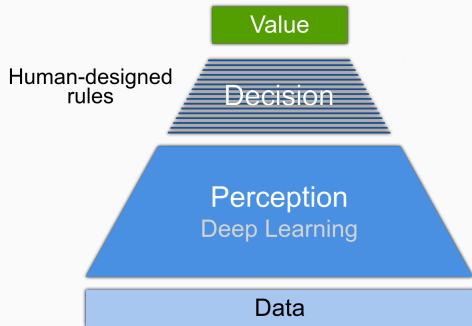
Problems involving an **agent** interacting with an **environment**, which provides numeric **reward signals**.

**Goal:** Learn how to take actions in order to maximize a reward function.



# From Data to Value

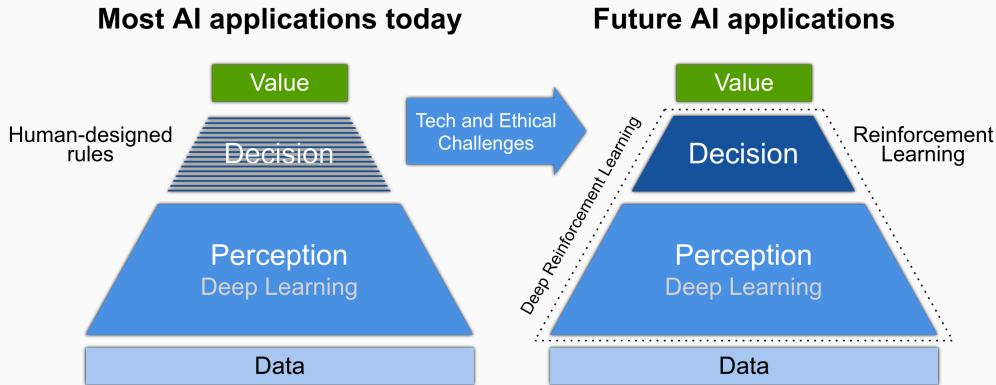
## Most AI applications today



---

Charafeddine, *Reinforcement Learning in the Wild and Lessons Learned*.

# From Data to Value



Charafeddine, *Reinforcement Learning in the Wild and Lessons Learned*.

# Components of the Agent

- **Policy:** agent's behaviour function

**Deterministic:**  $\pi(s) = a$

**Stochastic:**  $\pi(a|s) = \mathbb{P}[A_t = a|S_t = s]$

- **Value Function:** policy evaluation function

**State Value:**  $V^\pi(s) = \mathbb{E} \left[ \sum_{t \geq 0} \gamma^k r_t | S_0 = s, \pi \right]$

**Action Value:**  $Q^\pi(s, a) = \mathbb{E} \left[ \sum_{t \geq 0} \gamma^k r_t | S_0 = s, a_0 = a, \pi \right]$

- **Model:** agent's representation of the environment

# Categorizing Reinforcement Learning agents

- **Value Based**

- No Policy (implicit)
- Value Function

- **Policy Based**

- Policy
- No value function

- **Actor Critic**

- Policy
- Value function

- **Model Free**

- Policy and/or value function
- No Model

- **Model Based**

- Policy and/or value function
- Model

# Categorizing Reinforcement Learning agents

- **Value Based**

- No Policy (implicit)
- Value Function

- **Policy Based**

- Policy
- No value function

- **Actor Critic**

- Policy
- Value function

- **Model Free**

- Policy and/or value function
- No Model

- **Model Based**

- Policy and/or value function
- Model


# Model-Free Actor Critic methods

## Critic Network

Estimates the value function. This could be the action value  $Q$  or state value  $V$ .

## Actor Network

Updates the policy distribution in the direction suggested by the Critic (such as with policy gradients).



# Deep Deterministic Policy Gradient (DDPG)

- Off-Policy:
  - **Experience Replay Memory** of  $(s_t, a_t, r_t, s_{t+1}, d_t)$  tuples
- Action space: Continuous
- Policy: **Deterministic**
- Exploration:
  - **Ornstein-Uhlenbeck** process noise
  - Noise regulation with  **$\epsilon$ -decay function**

Needs accurate hyper-parameters fine-tuning

---

Lillicrap et al., “Continuous control with deep reinforcement learning”.



# Deep Deterministic Policy Gradient (DDPG) - Neural Networks

It uses **Target Networks** to minimise the instability MSBE loss

## 2 Local Neural Networks:

- Actor  $\pi(s \mid \theta)$
- Critic  $Q(s, a \mid \phi)$

## 2 Target Neural Networks:

- Actor  $\pi'(s \mid \bar{\theta})$
- Critic  $Q'(s, a \mid \bar{\phi})$

## Deep Deterministic Policy Gradient (DDPG) - Learning Equations

$$\begin{aligned} L(\phi) &= \mathbb{E}_{s_t \sim \rho^\beta, a_t \sim \beta, r_t \sim E} [(Q(s_t, a_t | \phi) - y_t)^2] \\ y_t &= r(s_t, a_t) + \gamma(1 - d_t)Q'(s_{t+1}, \pi'(s_t + 1 | \bar{\theta}) | \bar{\phi}) \end{aligned} \tag{1}$$

---

Lillicrap et al., “Continuous control with deep reinforcement learning”.

# Soft Actor-Critic (SAC)

- Off-Policy
- Action space: Continuous
- Policy: **Stochastic**
- Exploration: **Temperature Parameter**
- SAC is an off-policy algorithm which exploits entropy-regularized reinforcement learning
- Auto-tune parameters: Less hyper-parameters, less tuning
- Suitable for Real-World Experiments

---

Haarnoja et al., “Soft actor-critic algorithms and applications”.

# Design of the control system

---

## Experimental methodology and results

---

## Conclusions and future work

---









Thank you!



### References

---

-  Charafeddine, Mohamad. *Reinforcement Learning in the Wild and Lessons Learned*. 2018. URL: <https://link.medium.com/SRUZ24Itx4> (visited on 10/26/2018).
-  GovTech: Government Technology. *Autonomous Vehicles: Coming to a Road Near You*. 2018. URL: <https://www.govtech.com/transportation/Autonomous-Vehicles-Coming-to-a-Road-Near-You.html>.
-  Haarnoja, Tuomas et al. "Soft actor-critic algorithms and applications". In: *arXiv preprint arXiv:1812.05905* (2018).

-  Lillicrap, Timothy P et al. “Continuous control with deep reinforcement learning”. In: *arXiv preprint arXiv:1509.02971* (2015).
-  Pavone, Marco. *Veicoli a guida autonoma: a che punto siamo e cosa ci aspetta?* Festival della Tecnologia Conference. 2019.
-  Sutton, Richard S and Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 2018.